INFOCEPTS
Data & AI

FutureCart CRM

Project Report

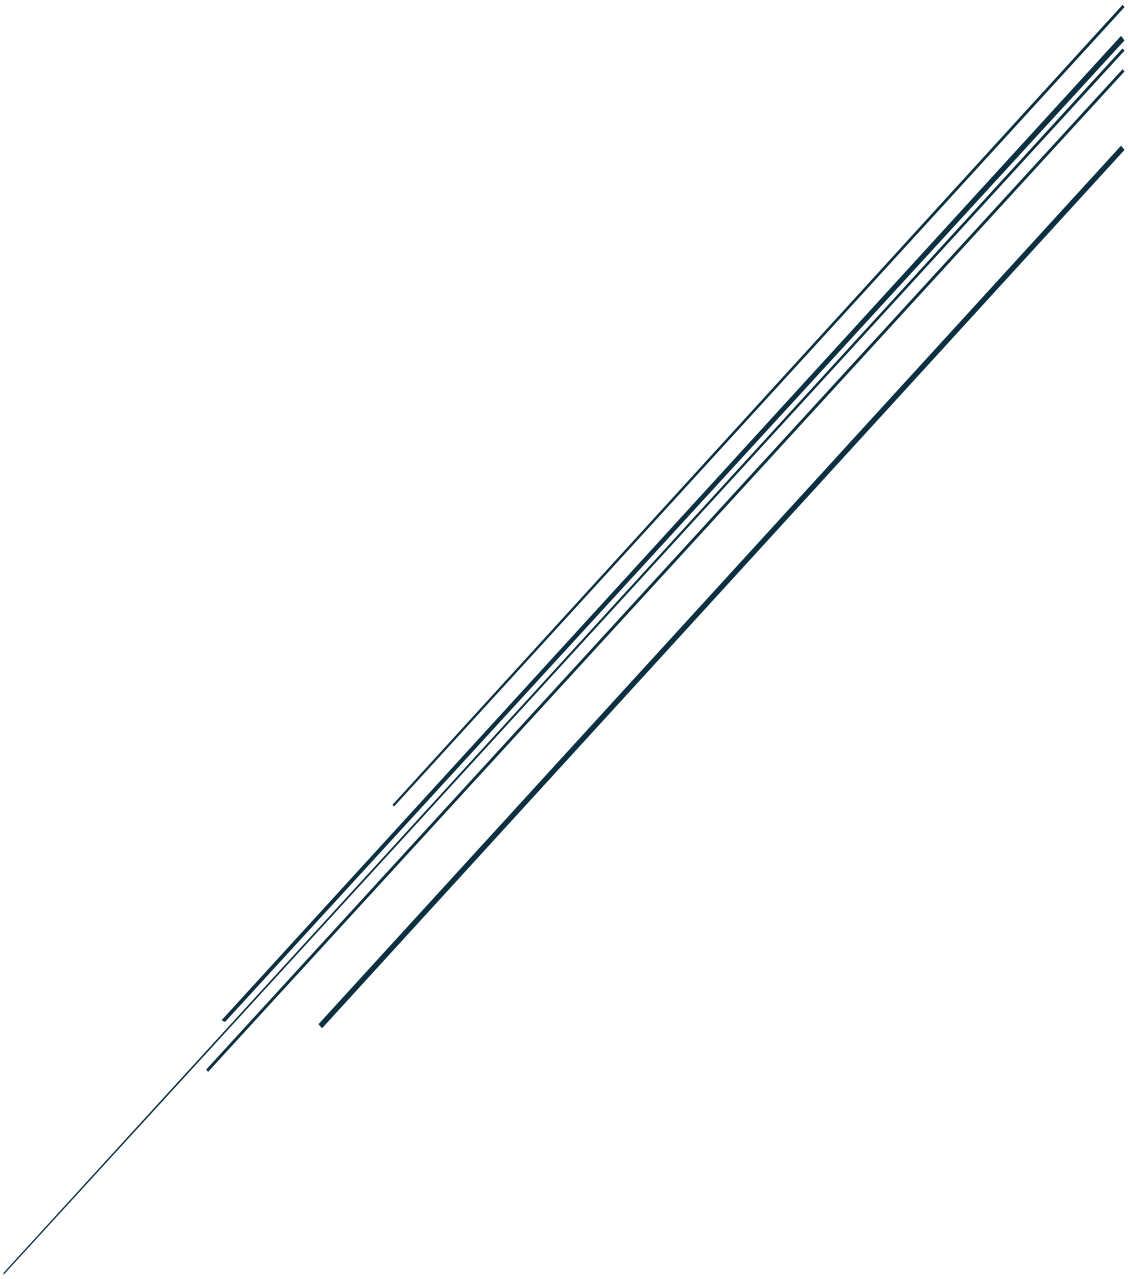

By
Aahash Kamble

Problem Statement

FutureCart Inc., a leading retail and e-commerce company in India, aims to enhance its customer retention and service quality by leveraging a modern CRM analytics platform. The company operates multiple contact centers and collects customer interaction data through various channels such as calls, chats, and emails. To transform this data into actionable insights, FutureCart needs a robust data infrastructure that supports both batch and real-time data processing.

The core challenge is to build a CRM Data Mart using a Lambda Architecture that can ingest, process, and analyze both historical and streaming data. This system should enable the generation of real-time and batch KPIs to support strategic decision-making and improve customer satisfaction.

Overview: FutureCart CRM – Lambda Architecture for Real-Time and Batch Analytics Flow Digram

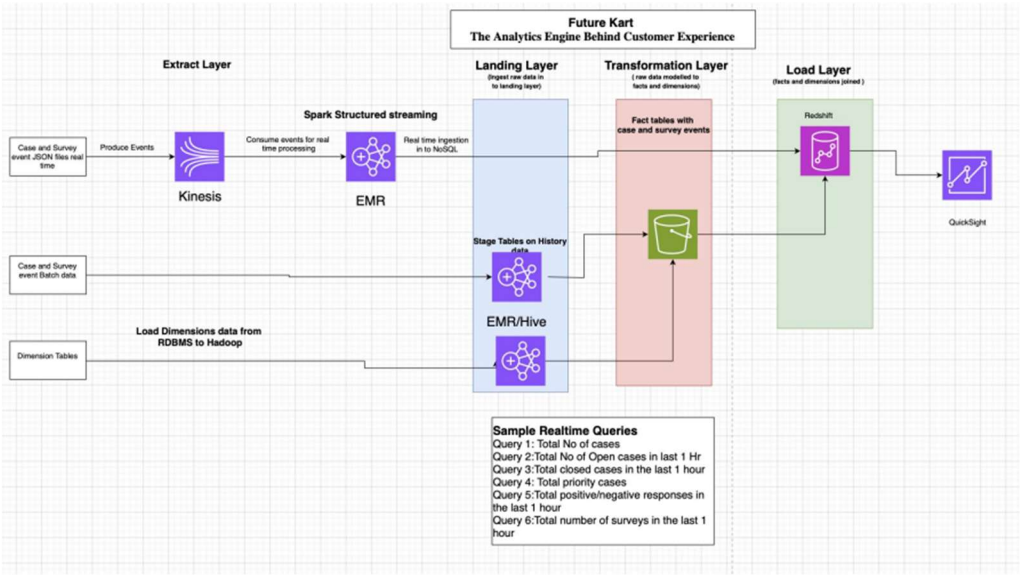

Key Technical Objectives

1. Set up MySQL on EC2 to store dimension data.
2. Ingest dimension data into Hive for batch processing.
3. Generate and load historical case and survey data into HDFS.
4. Stream real-time events (cases and surveys) using AWS Kinesis.
5. Process real-time data using Spark Structured Streaming and load it into Amazon Redshift.
6. Export historical and dimension data from Hive to S3, then load into Redshift.
7. Run analytical queries on Redshift to compute KPIs.

Phase 1: Infrastructure Setup & Initial Configuration

Tasks:

- Launched **EC2 instances** for local processing.
- Installed **MySQL 5.7** and configured schema.
- Listed all S3 buckets and verified access to aahash-project.
- Created schema and tables in **MySQL** for dimension data.

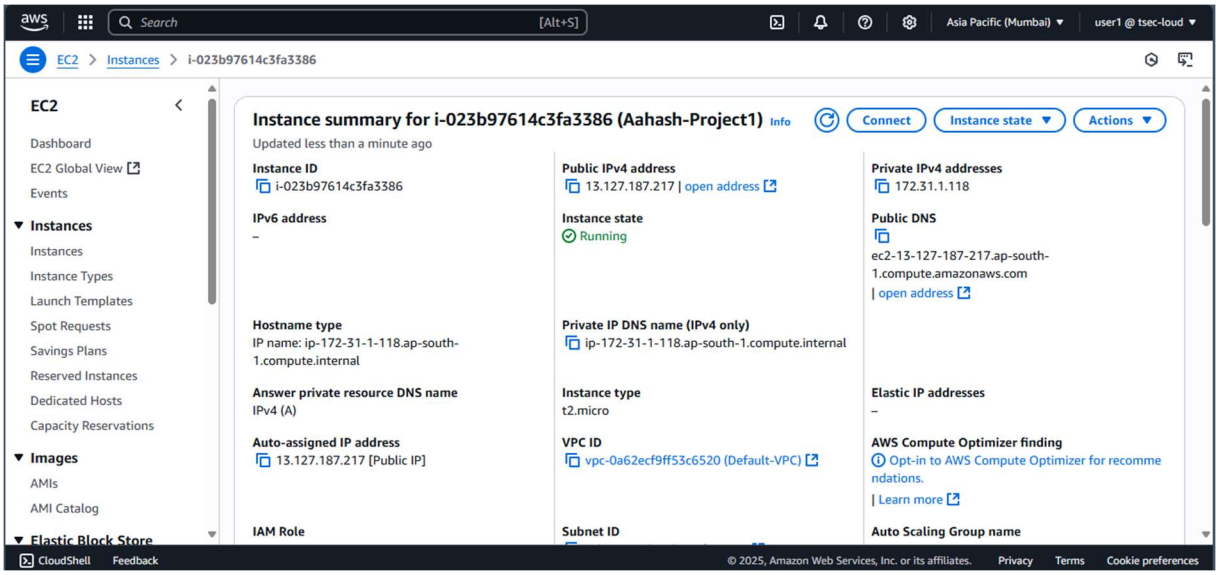

Install and setup Mysql
sudo yum update -y
sudo yum install -y mariadb-server
sudo systemctl start mariadb
sudo systemctl enable mariadb

Create MySQL Tables and Load Data

CREATE TABLE calendar_details (calendar_date DATE, date_desc VARCHAR(50), week_day_nbr SMALLINT, week_number SMALLINT, week_name VARCHAR(50), year_week_number INT, month_number SMALLINT, month_name VARCHAR(50), quarter_number SMALLINT, quarter_name VARCHAR(50), half_year_number SMALLINT, half_year_name VARCHAR(50), geo_region_cd CHAR(2));

LOAD DATA LOCAL INFILE '/home/ec2-user/dimensions/futurecart_calendar_details.txt' INTO TABLE calendar_details FIELDS TERMINATED BY ',' ENCLOSED BY '"' LINES TERMINATED BY '\n' IGNORE 1 LINES;

Tables are created :

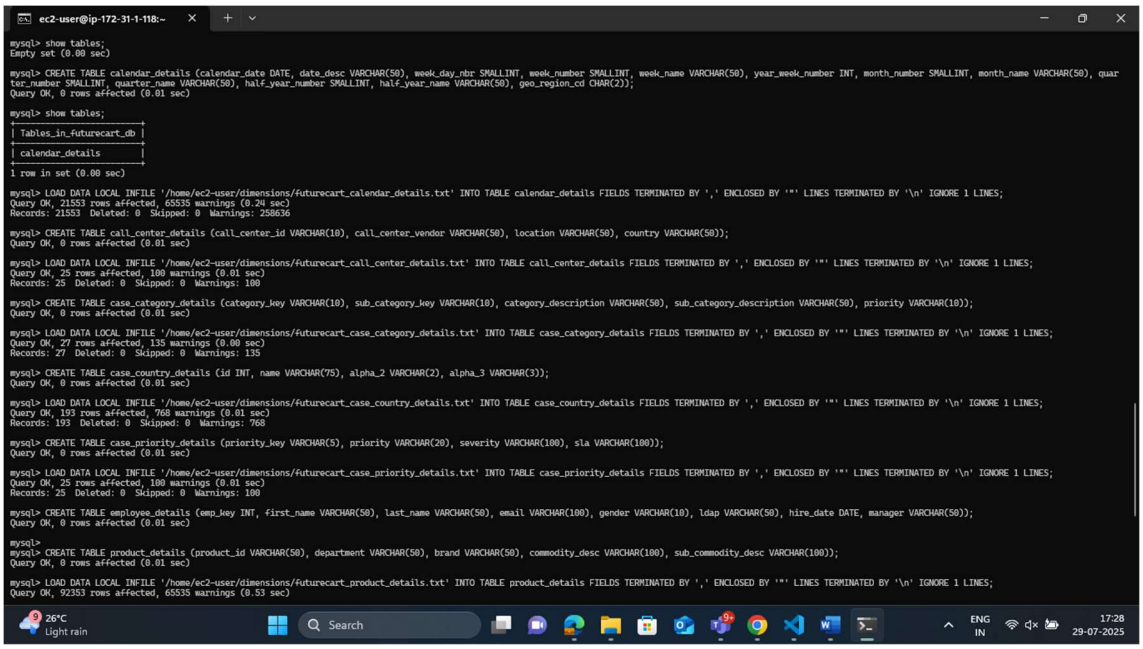


```
mysql> show tables;
+--------------------------+
| Tables_in_futurecart_db  |
+--------------------------+
| calendar_details         |
| call_center_details      |
| case_category_details    |
| case_country_details     |
| case_priority_details    |
| employee_details         |
| product_details          |
| survey_question_details  |
+--------------------------+
8 rows in set (0.00 sec)

mysql>
```

Phase 2: Dimension Data Preparation & Upload

Tasks:

- Downloaded *.txt files from s3://aahash-project/dimensions/.
- Uploaded these files to EC2 instance in ~/futurecart_dims/.
- Loaded these files into **HDFS** via Hadoop command-line.
- Created **Hive tables** and loaded the data into Hive.
- Exported Hive tables as **JSON** files.
- Uploaded the JSON exports to s3://aahash-project/dimensions/.

Load Data to Hadoop (EMR)

```
aws s3 cp s3://aahash-project/dimensions/ ~/futurecart_dims/ --recursive
hadoop fs -mkdir -p /user/ec2-user/futurecart_dims
hadoop fs -put ~/futurecart_dims/*.txt /user/ec2-user/futurecart_dims/

Create Hive external tables
CREATE EXTERNAL TABLE calendar_details (
  calendar_date STRING,
  date_desc STRING,
  week_day_nbr INT,
  week_number INT,
  week_name STRING,
  year_week_number STRING,
  month_number INT,
  month_name STRING,
  quarter_number INT,
  quarter_name STRING,
  half_year_number INT,
  half_year_name STRING,
  geo_region_cd STRING
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE

LOCATION 'hdfs:///user/ec2-user/futurecart_dims/calendar_details';
```

```
[ec2-user@ip-172-31-22-83 ~]$ hadoop fs -mkdir -p /user/ec2-user/futurecart_dims
[ec2-user@ip-172-31-22-83 ~]$ hadoop fs -put ~/futurecart_dims/*.txt /user/ec2-user/futurecart_dims/
put: '/home/ec2-user/futurecart_dims/*.txt': No such file or directory
[ec2-user@ip-172-31-22-83 ~]$ hive
Hive Session ID = f51ef723-e107-4d1d-a6fb-28c3e72a749a

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: true
hive> CREATE EXTERNAL TABLE calendar_details (
    >   calendar_date STRING,
    >   date_desc STRING,
    >   week_day_nbr INT,
    >   week_number INT,
    >   week_name STRING,
    >   year_week_number STRING,
    >   month_number INT,
    >   month_name STRING,
    >   quarter_number INT,
    >   quarter_name STRING,
    >   half_year_number INT,
    >   half_year_name STRING,
    >   geo_region_cd STRING
    > )
    > ROW FORMAT DELIMITED
    > FIELDS TERMINATED BY '\t'
    > STORED AS TEXTFILE
    > LOCATION 'hdfs:///user/ec2-user/futurecart_dims/calendar_details';
OK
Time taken: 1.957 seconds
hive> CREATE TABLE calendar_json
    > ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe'
    > STORED AS TEXTFILE
    > LOCATION '/user/ec2-user/futurecart_json/calendar/'
    > AS SELECT * FROM calendar_details;
Query ID = ec2-user_20250731083315_dc5c76c4-c28b-4e2b-8dd6-52e91e637f3c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1753936989620_0003)
```

```
Time taken: 1.957 seconds
hive> CREATE TABLE calendar_json
    > ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe'
    > STORED AS TEXTFILE
    > LOCATION '/user/ec2-user/futurecart_json/calendar/'
    > AS SELECT * FROM calendar_details;
Query ID = ec2-user_20250731083315_dc5c76c4-c28b-4e2b-8dd6-52e91e637f3c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1753936989620_

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FA
--------------------------------------------------------------------------------
Map 1            container     SUCCEEDED      0          0        0        0
--------------------------------------------------------------------------------
VERTICES: 00/01  [>>--------------------------] 0%    ELAPSED TIME: 0.17 s
--------------------------------------------------------------------------------
Moving data to directory /user/ec2-user/futurecart_json/calendar
OK
Time taken: 3.243 seconds
hive> hadoop fs -get /user/ec2-user/futurecart_json/calendar/ ~/calendar_json/
```

Copy JSON output from HDFS to local:

hadoop fs -get /user/ec2-user/futurecart_json/calendar/ ~/calendar_json/

Upload to S3:

aws s3 cp ~/calendar_json/ s3://aahash-project/dimensions/calendar_json/ --recursive

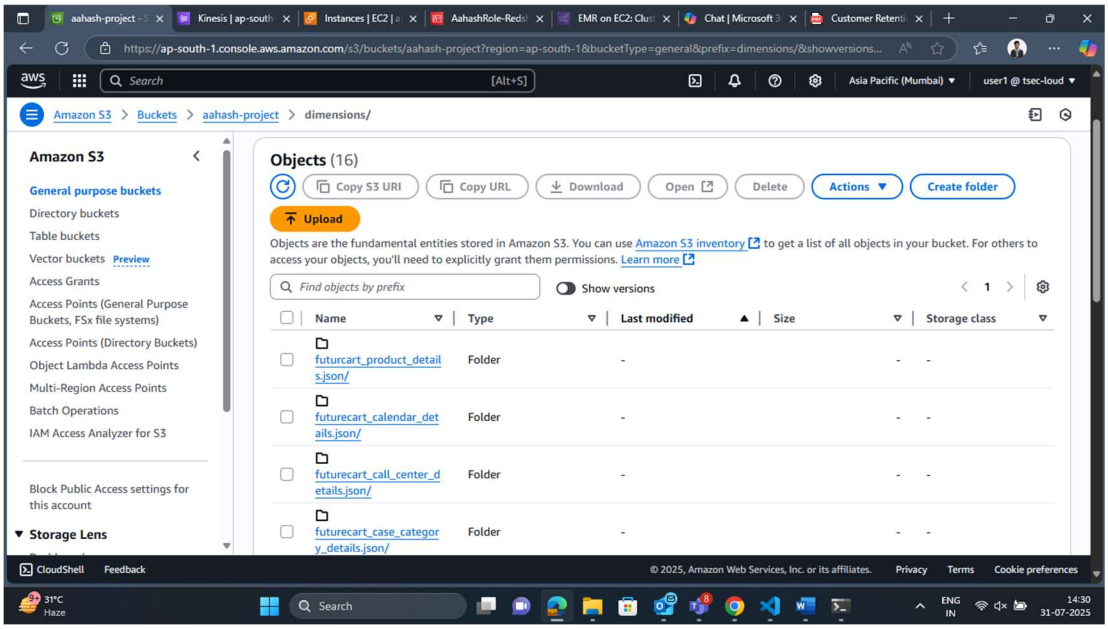

Phase 3: Historical Data Generation and Batch Processing

Tasks:

- Used a Python script to generate historical **case/survey** data in JSON.
- Uploaded the generated data to s3://aahash-project/historical-data/.

mkdir -p ~/scripts
copy scripts
aws s3 cp s3://aahash-project/scripts/ ~/scripts/ --recursive
cd ~/scripts
python3 generate_historical_data.py
aws s3 cp s3://aahash-project/historical_data/ ~/historical_data/ --recursive

Script to generate historical data

```
import os
import json
import random
from datetime import datetime, timedelta

# CONFIGURABLE PARAMETERS
NUM_DAYS = 10
CASES_PER_DAY = 100
SURVEYS_PER_DAY = 80  # must be ≤ CASES_PER_DAY
OUTPUT_DIR = "historical_data_jsonl"
def random_case_no():
    return str(random.randint(600000, 700000))
def random_timestamp(day_offset):
    date = datetime.now() - timedelta(days=day_offset)
    return date.strftime("%Y-%m-%d %H:%M:%S")
def generate_case(case_no, day_offset):
    return {
        "status": random.choice(["Open", "Closed"]),
        "category": f"CAT{random.randint(1,5)}",
        "sub_category": f"SCAT{random.randint(1,20)}",
        "last_modified_timestamp": random_timestamp(day_offset),
        "case_no": case_no,
        "create_timestamp": random_timestamp(day_offset),
        "created_employee_key": str(random.randint(200000, 300000)),
        "call_center_id": f"C-{random.randint(100,120)}",
        "product_code": str(random.randint(9000000, 9999999)),
        "country_cd": random.choice(["IN", "US", "BR", "DE", "AU"]),
        "communication_mode": random.choice(["Email", "Call", "Chat"])
    }

def generate_survey(case_no, day_offset):
    return {
        "survey_id": f"S-{random.randint(500000, 599999)}",
```

```
        "case_no": case_no,
        "survey_timestamp": random_timestamp(day_offset),
        "Q1": random.randint(1, 10),
        "Q2": random.randint(1, 10),
        "Q3": random.randint(1, 10),
        "Q4": random.choice(["Y", "N"]),
        "Q5": random.randint(1, 10)
    }

def ensure_dirs():
    os.makedirs(f"{OUTPUT_DIR}/cases", exist_ok=True)
    os.makedirs(f"{OUTPUT_DIR}/surveys", exist_ok=True)

def main():
    ensure_dirs()
    for day in range(1, NUM_DAYS + 1):
        case_nos = [random_case_no() for _ in range(CASES_PER_DAY)]
        cases = [generate_case(cn, day) for cn in case_nos]
        surveys = [generate_survey(cn, day) for cn in random.sample(case_nos, SURVEYS_PER_DAY)]

        # Write case data as JSON Lines
        with open(f"{OUTPUT_DIR}/cases/case_data_day{day}.json", "w") as f:
            for record in cases:
                f.write(json.dumps(record) + "\n")

        # Write survey data as JSON Lines
        with open(f"{OUTPUT_DIR}/surveys/survey_data_day{day}.json", "w") as f:
            for record in surveys:
                f.write(json.dumps(record) + "\n")

    print(f"✅ JSON Lines generated for {NUM_DAYS} days in '{OUTPUT_DIR}/'")

if __name__ == "__main__":
    main()
```

Loaded and generated Json case event

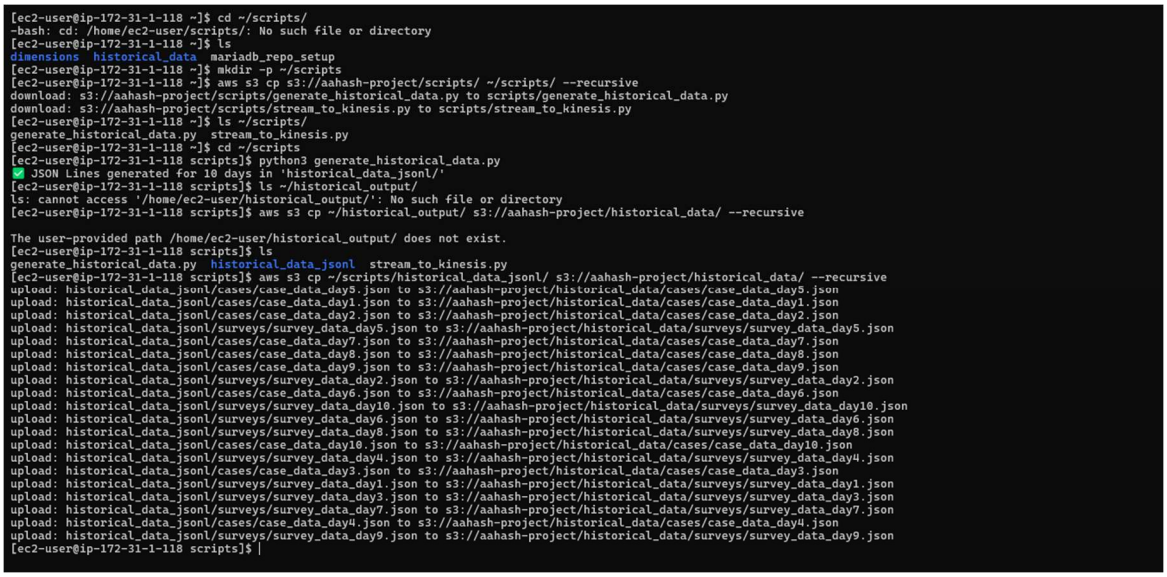

Phase 4: Real-Time Data Ingestion via Kinesis

Tasks:

- Created a **Kinesis Data Stream**: aahash-newstream.
- Built send_to_kinesis.py to push real-time events to the stream in JSON format.
- Validated Kinesis streaming using logs and AWS Console.

Send data to kinesis

STREAM_NAME = "aahash-data"

python3 stream_to_kinesis.py

```
import boto3
import random
import time
import calendar
import json
from datetime import datetime, timedelta

# Initialize Kinesis client
kinesis = boto3.client('kinesis', region_name='ap-south-1')
stream_name = 'aahash-stream'

# Read case data
with open('realtimedata/000000_0', 'r') as case_data_obj:
    all_case_data_lines = case_data_obj.readlines()

open_case_time_diff_mins = [40, 50, 60]
closed_case_time_diff_mins = [5, 10, 20, 30]
number_of_cases_counts = [1, 2, 3, 4, 5, 6]
scores = list(range(1, 11))
answer = ["Y", "N"]
survey_id_start = 500000
category = 'CAT3'
sub_categorys = ['SCAT8', 'SCAT9', 'SCAT10', 'SCAT11', 'SCAT12', 'SCAT13', 'SCAT14', 'SCAT15', 'SCAT16']

total_cases = len(all_case_data_lines)
i = 900

while i <= (total_cases - 1):
    sub_category = random.choice(sub_categorys)
    number_of_cases = random.choice(number_of_cases_counts)
    cases = all_case_data_lines[i:i + number_of_cases]

    current_timestamp = datetime.now()
    case_created_ts = str(current_timestamp - timedelta(minutes=random.choice(open_case_time_diff_mins)))[:19]
    case_closed_ts = str(current_timestamp - timedelta(minutes=random.choice(closed_case_time_diff_mins)))[:19]
    survey_ts = str(current_timestamp)[:19]
    file_ts = calendar.timegm(time.gmtime())

    for j in cases:
        case_no, created_employee, call_center, status, category1, sub_category1, mode, country, product = j.strip().split(',')

        case_data = {
            "case_no": case_no,
            "created_employee_key": created_employee,
```

```
            "call_center_id": call_center,
            "status": "Closed" if i % 5 == 0 else status,
            "category": category,
            "sub_category": sub_category,
            "communication_mode": mode,
            "country_cd": country,
            "product_code": product,
            "last_modified_timestamp": case_closed_ts if i % 5 == 0 else case_created_ts,
            "create_timestamp": case_created_ts
        }

        # Send to Kinesis
        kinesis.put_record(
            StreamName=stream_name,
            Data=json.dumps(case_data),
            PartitionKey=case_no
        )
        print("Sent to Kinesis:", case_data)

        if i % 5 == 0:
            survey_data = {
                "survey_id": f"S-{survey_id_start}",
                "case_no": case_no,
                "survey_timestamp": survey_ts,
                "Q1": random.choice(scores),
                "Q2": random.choice(scores),
                "Q3": random.choice(scores),
                "Q4": random.choice(answer),
                "Q5": random.choice(scores)
            }
            kinesis.put_record(
                StreamName=stream_name,
                Data=json.dumps(survey_data),
                PartitionKey=case_no
            )
            print("Sent survey to Kinesis:", survey_data)
            survey_id_start += 1

    i += number_of_cases
    time.sleep(5)
```

Verify stream data :

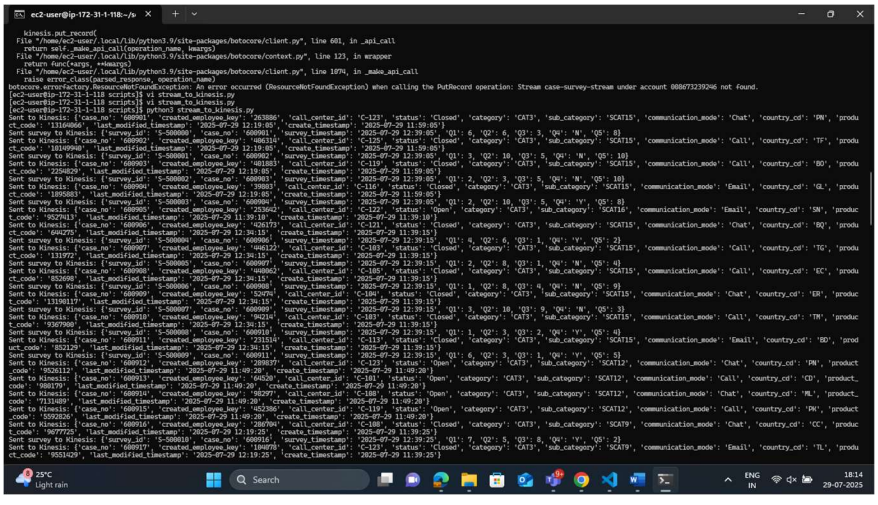

Phase 5: Real-Time Stream Processing with Spark

Tasks:

- Developed kinesis_to_redshift.py using Spark Structured Streaming.
- Set up Spark to ingest real-time JSON events from Kinesis stream (aahash-newstream).
- Transformed and batched incoming data for compatibility with Amazon Redshift schema.
- Wrote transformed data into Redshift tables using JDBC connector.

Created Redshift

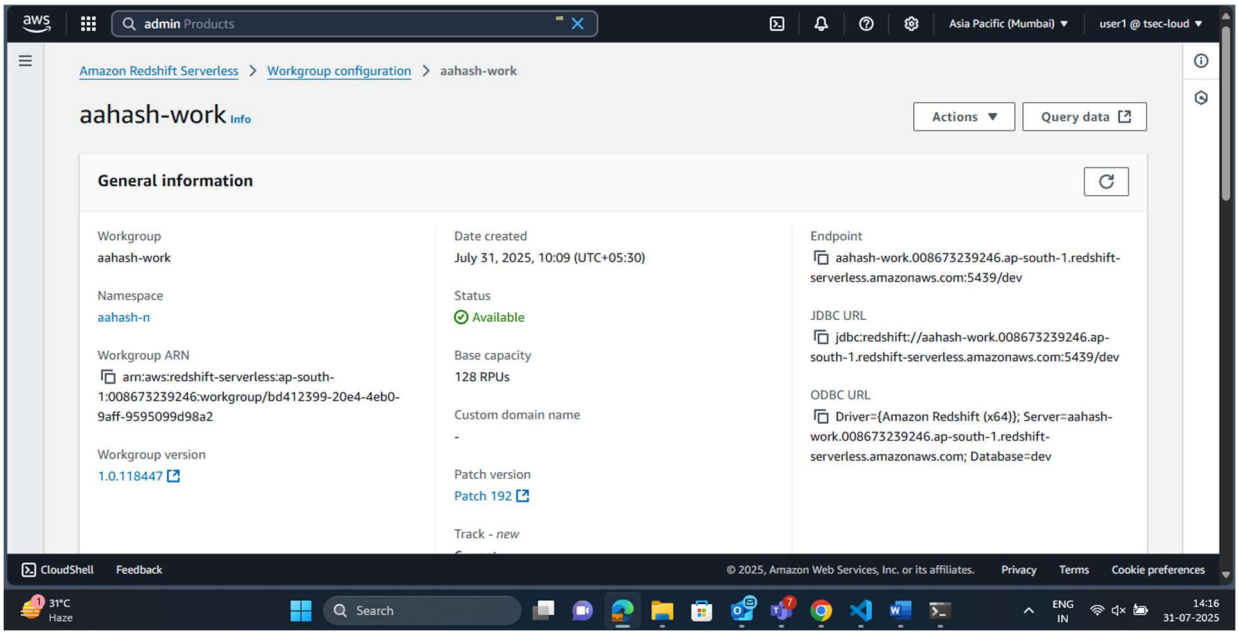

Run the job using spark-submit on the EMR cluster:

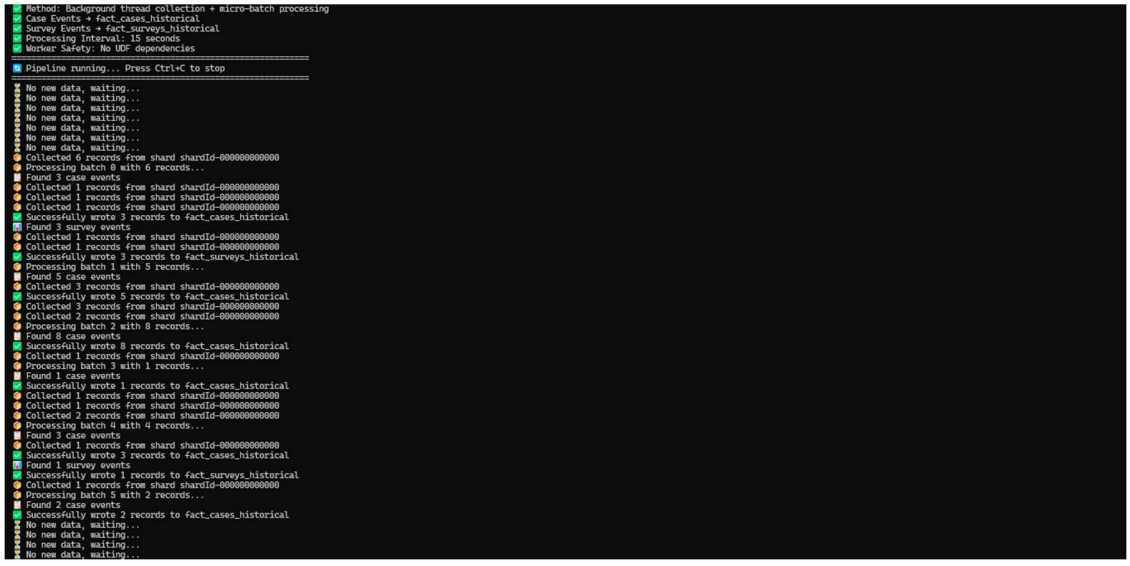

Data is received by Redshift:

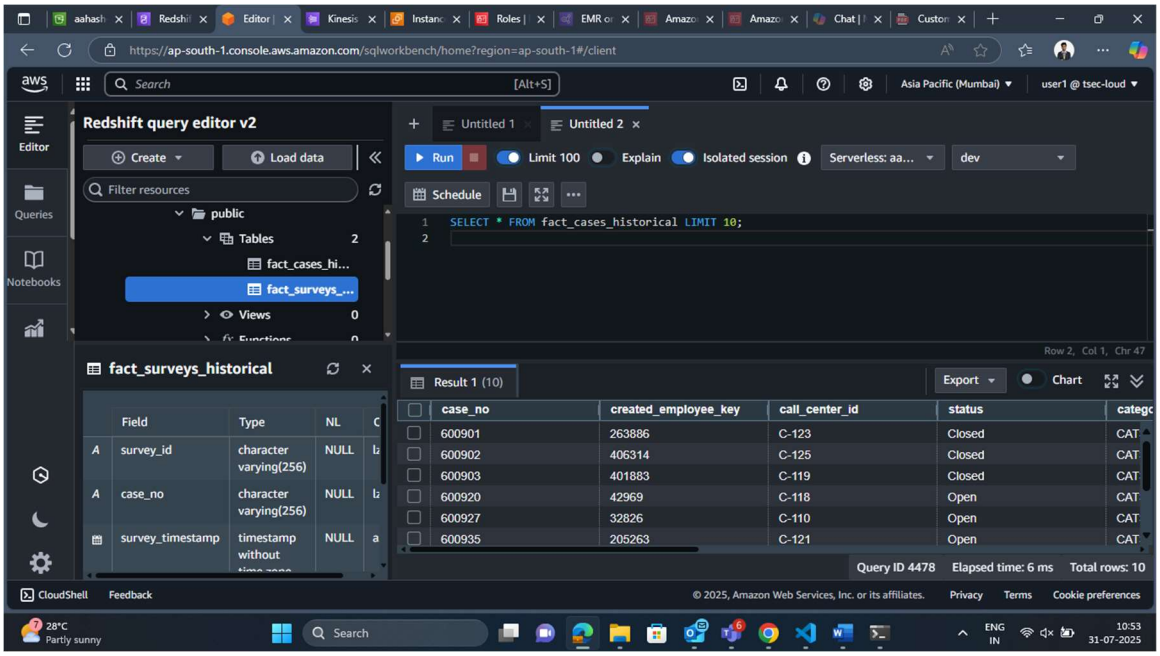

Phase 6: Query Execution and Final Insights

Tasks:

- Loaded dimension JSON data from S3 to local SQL.
- Verified integrity of data using SELECT queries.
- Ran provided analytical queries on the complete dataset.

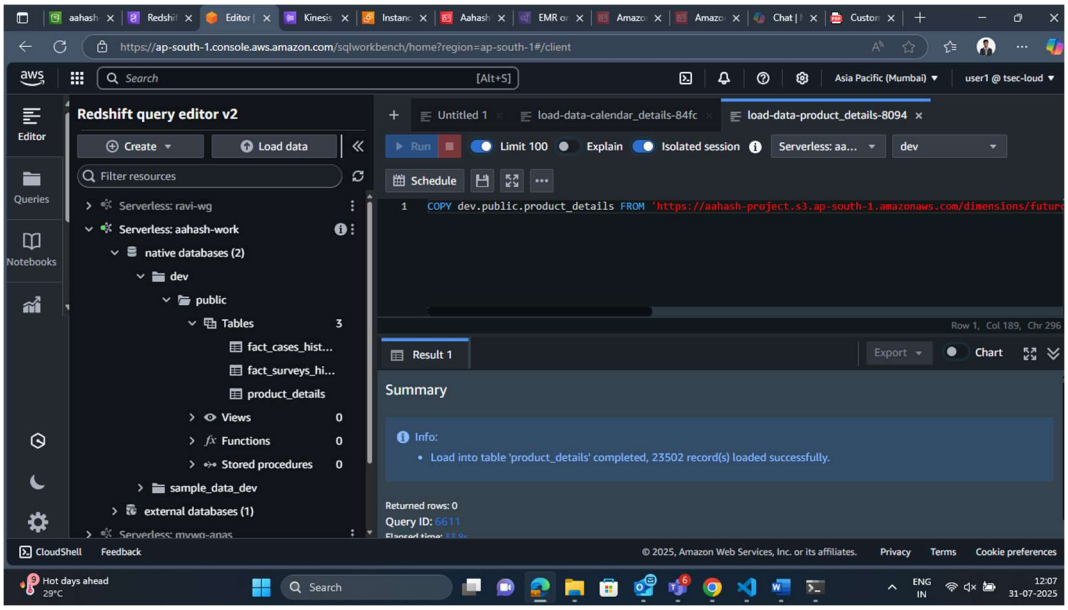

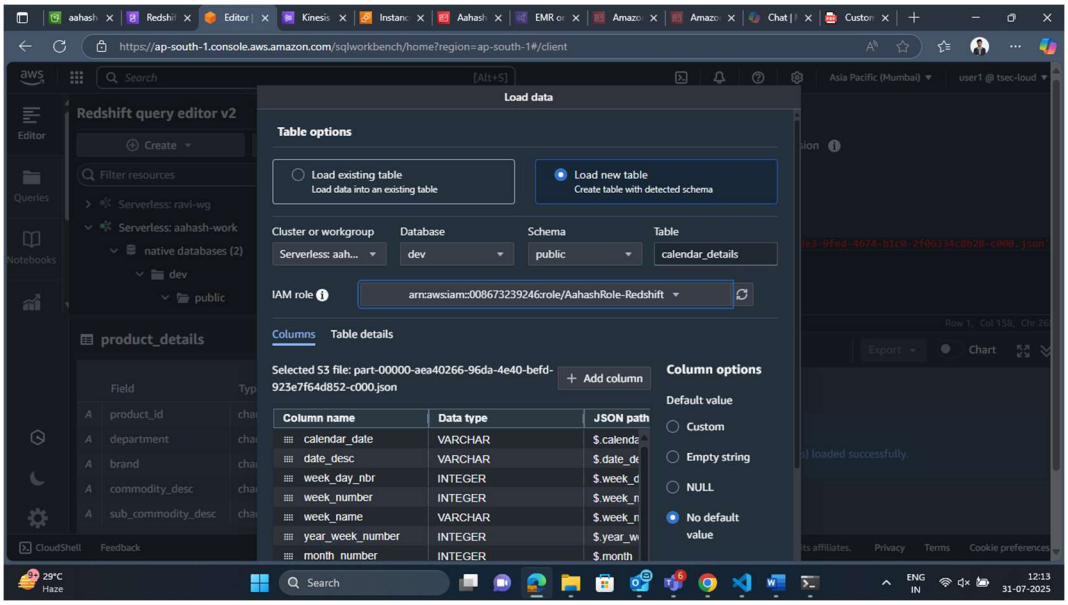

All data has been loaded:

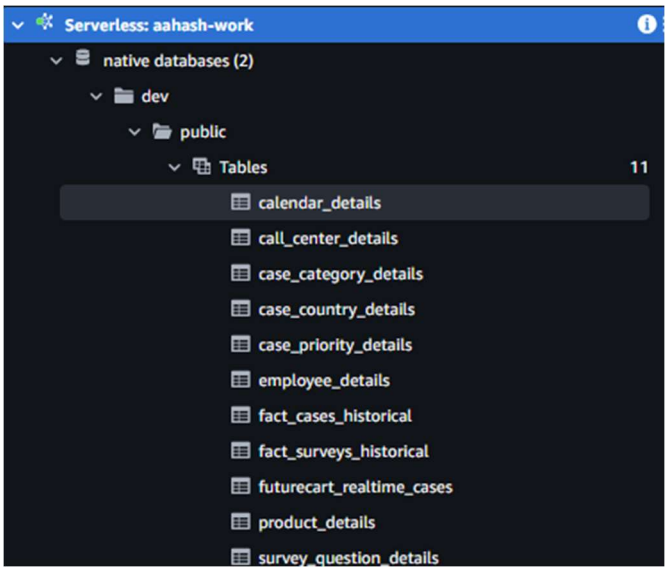

Create the above queries on the Redshift tables

1. Total Number of Cases

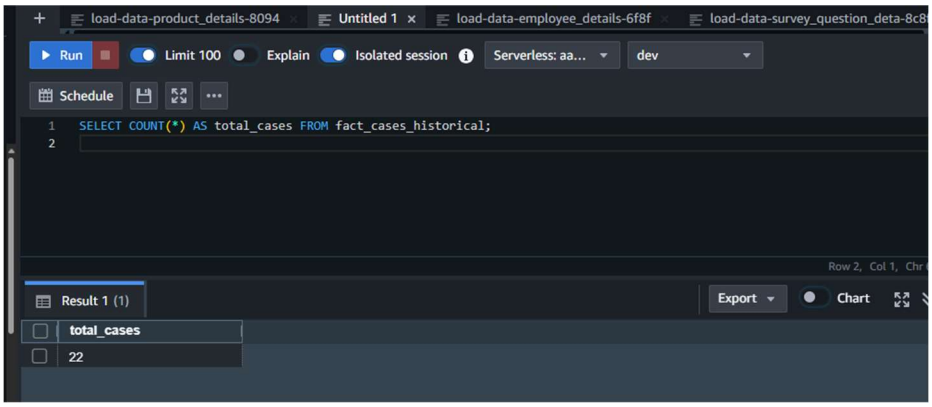

2. Total Open/Closed Cases in the Last 1 Hour

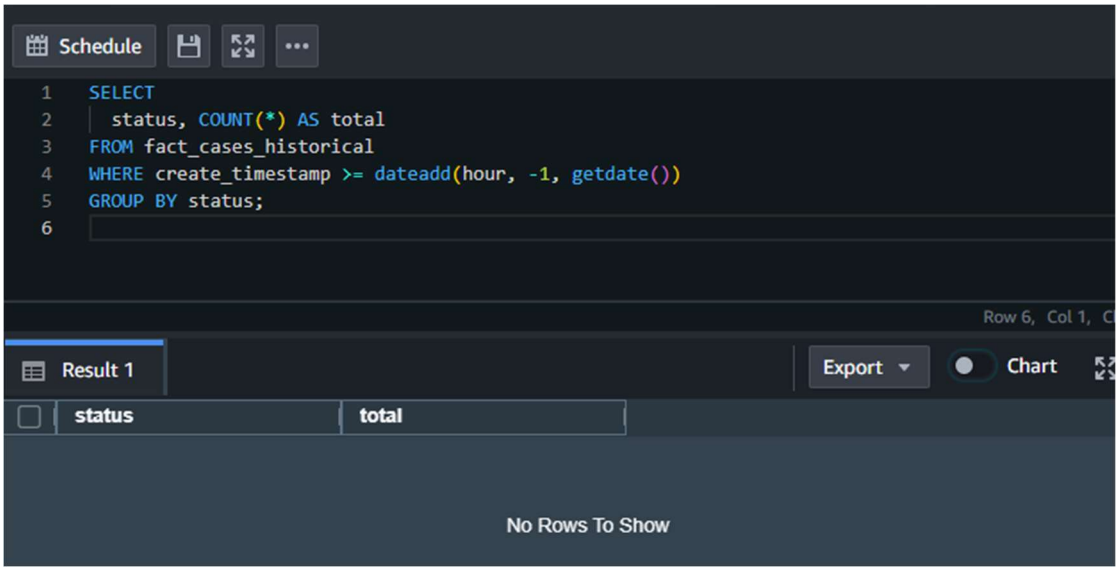

3. Total Priority Cases (Real-Time)

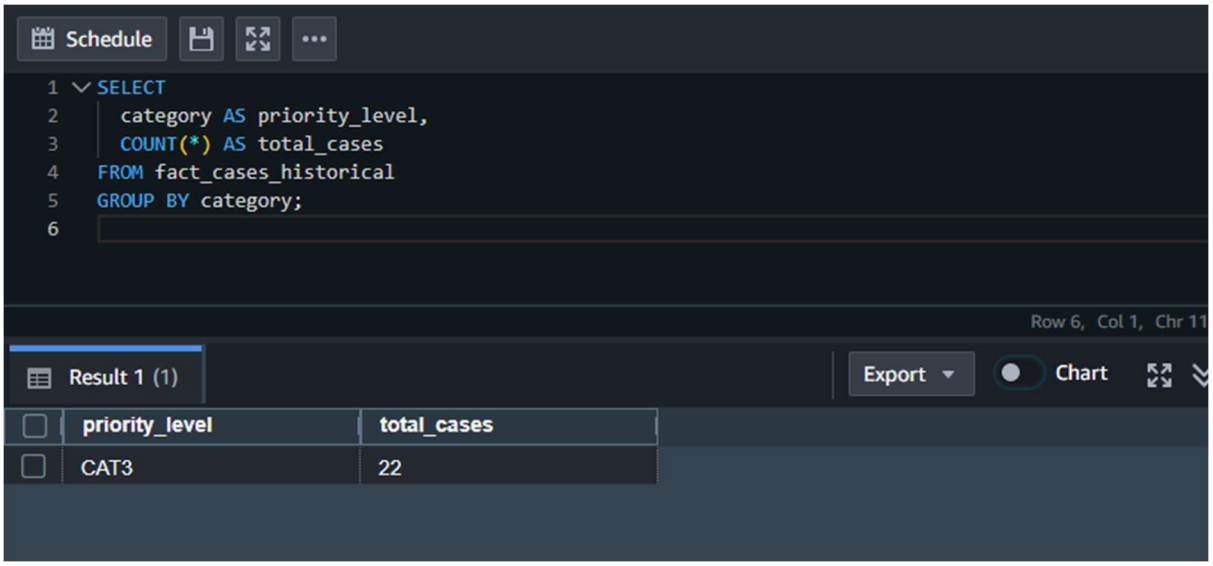

4. Total Positive/Negative Responses in the Last 1 Hour (Real-Time)

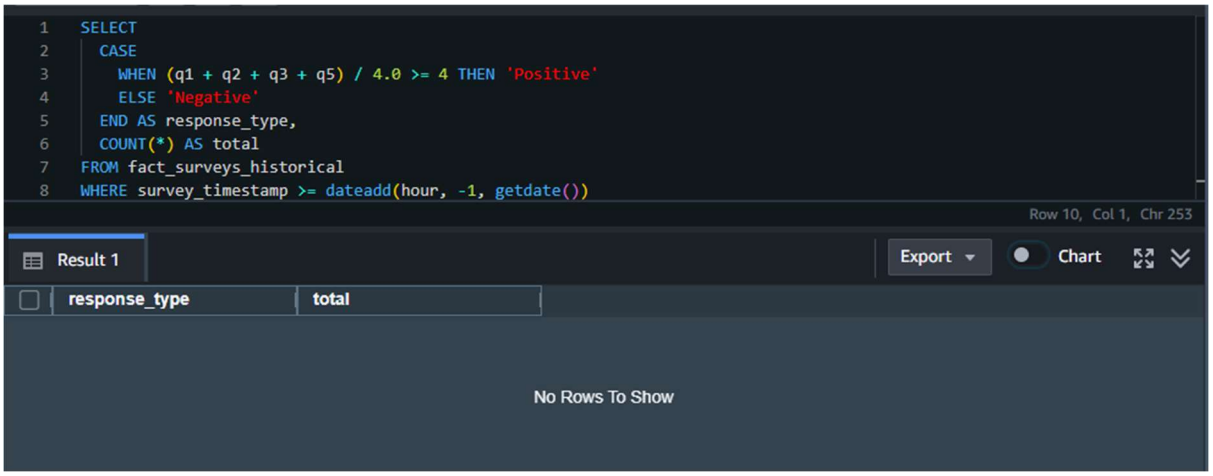

5: Total Number of Surveys in the Last 1 Hour (Real-Time)

6: Total Open/Closed Cases in a Day/Week/Month (Historical)

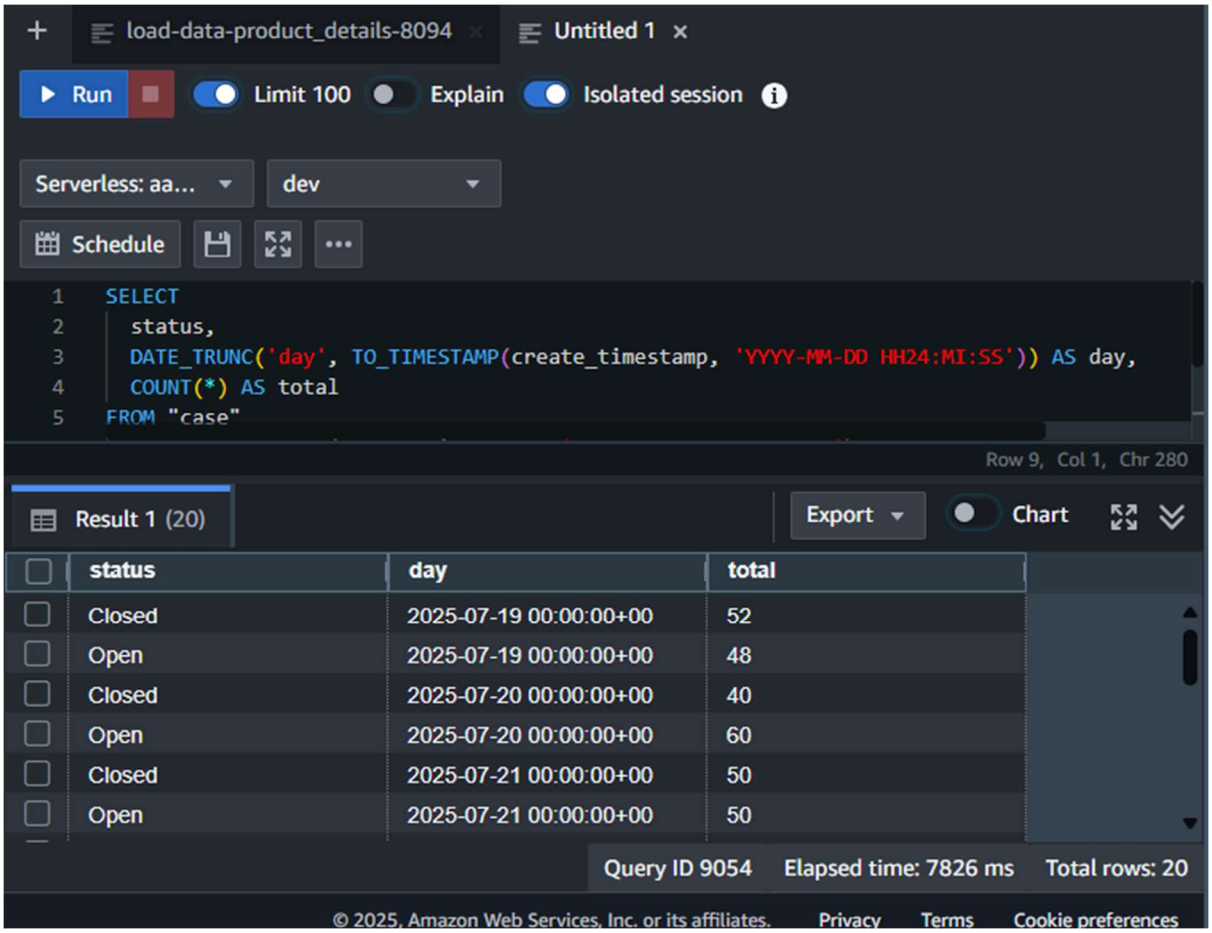

7: Total Positive/Negative Responses in a Day/Week/Month (Historical)

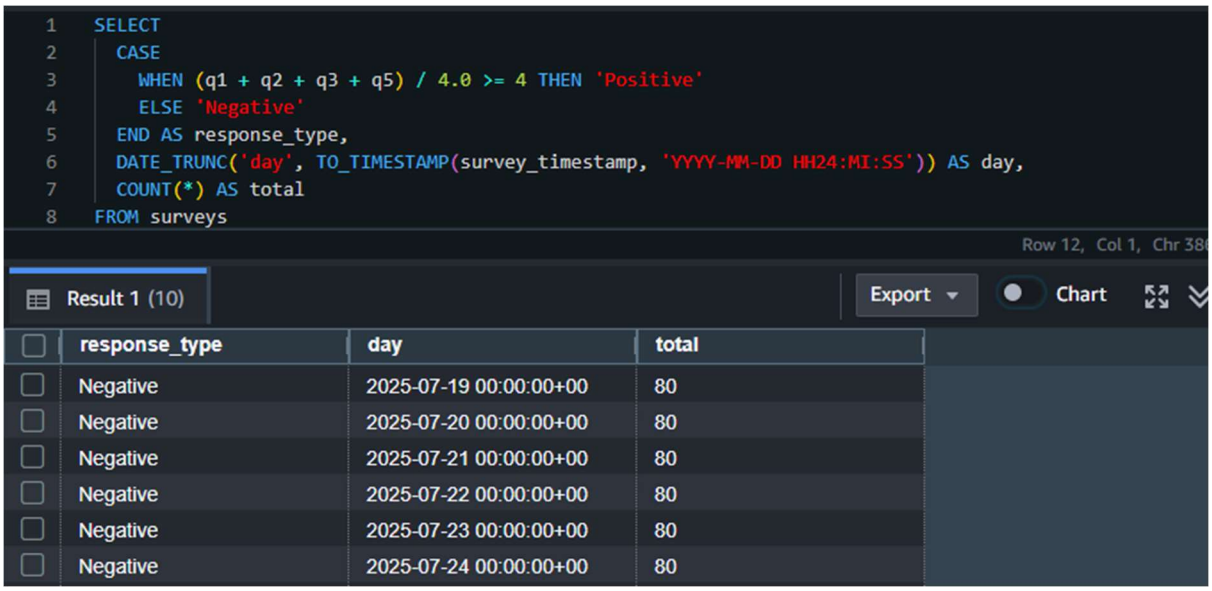

8: Total Number of Surveys in a Day/Week/Month (Historical)

```
SELECT
  DATE_TRUNC('day', TO_TIMESTAMP(survey_timestamp, 'YYYY-MM-DD HH24:MI:SS')) AS day,
  COUNT(*) AS total_surveys
FROM surveys
WHERE TO_TIMESTAMP(survey_timestamp, 'YYYY-MM-DD HH24:MI:SS') >= CURRENT_DATE - INTERVAL '1 month'
GROUP BY day
ORDER BY day;
```

Row 8, Col 1, Ch

Result 1 (10) | Export | Chart

day	total_surveys
2025-07-19 00:00:00+00	80
2025-07-20 00:00:00+00	80
2025-07-21 00:00:00+00	80
2025-07-22 00:00:00+00	80
2025-07-23 00:00:00+00	80
2025-07-24 00:00:00+00	80

9: Real-time KPIs – Open vs Closed Cases (from Real-Time Table)

```
SELECT
  status, COUNT(*) AS total
FROM fact_cases_historical
GROUP BY status;
```

Result 1 (2)

status	total
Closed	4
Open	18

10: Cases Received by Priority and Severity (Real-Time)

```
SELECT
  category AS priority,
  sub_category AS severity,
  COUNT(*) AS total
```

Row 7, Col 1, Chr

Result 1 (8) | Export | Chart

priority	severity	total
CAT3	SCAT10	6
CAT3	SCAT9	3
CAT3	SCAT15	2
CAT3	SCAT11	1
CAT3	SCAT8	4
CAT3	SCAT13	4
CAT3	SCAT14	1
CAT3	SCAT16	1